
*QUEEN DOMINATION VIA SAT SOLVING +

Introduction

Given an $n \times n$ chessboard, the queen domination problem asks for the minimum number of queens required to dominate all squares of the board. A queen dominates every square lying in the same row, column, diagonal, or anti-diagonal as its position. The corresponding minimum number of queens is denoted by $\gamma(Q_n)$. Several theoretical results provide useful bounds on $\gamma(Q_n)$ and substantially reduce the search space. In particular, the inequality $\gamma(Q_n) - \gamma(Q_{n-1}) \leq 1$ limits the rate at which the domination number can grow, while the lower bound $\gamma(Q_n) \geq \lceil (n-1)/2 \rceil$ further restricts the range of feasible values. Together, these bounds often reduce the number of candidate values for $\gamma(Q_n)$ to only one or two for a given board size.

Despite these results, determining the exact domination number or enumerating all optimal solutions up to isomorphism generally requires computational methods. Most approaches proposed in the literature rely on highly specialized search procedures designed specifically for the combinatorial structure of the problem. More recently, however, it was shown that carefully engineered SAT encodings can achieve performance competitive with specialized solvers for queen domination [1]. In particular, the SAT-based encoding of [1] combined literal ordering heuristics, symmetry breaking constraints, and Cube-and-Conquer techniques to obtain strong computational performance, correct previously reported erroneous results, and resolve a previously open instance in the model enumeration setting.

Motivated by these developments, this report investigates a new SAT formulation for the queen domination problem that significantly outperforms the previous SAT-based approach. The formulation extends the encoding of [1] by introducing auxiliary variables associated with board lines, namely rows, columns, diagonals, and anti-diagonals. These variables expose additional geometric structure to the SAT solver and enable stronger forms of global reasoning during search.

The central observation underlying the formulation is straightforward: each queen activates exactly one row, one column, one diagonal, and one anti-diagonal. Consequently, any placement using at most γ queens can activate at most 4γ lines in total. This makes it possible to impose an additional cardinality constraint over the line variables, which in practice turns out to be one of the primary sources of the observed performance improvement.

The encoding further incorporates several optimizations, including Hilbert curve variable ordering, static symmetry breaking constraints, and a Cube-and-Conquer strategy tailored to totalizer-based cardinality encodings. In the model enumeration setting, previous work first enumerated all solutions and then performed a separate final UNSAT verification step [1]. In contrast, the present approach employs `cadical-exhaust`² together with `drat-trim-t`³, allowing solving and proof verification to be performed simultaneously.

Encoding Queen Domination with Line Variables

We consider the queen graph of an $n \times n$ chessboard, where each square corresponds to a vertex and edges connect squares that share a row, column, diagonal, or anti-diagonal. The objective is to determine whether there exists a dominating set of size at most γ .

Our formulation extends the encoding introduced in previous work [1] by introducing auxiliary variables corresponding to board lines. This additional structure enables the SAT solver to reason more globally about the geometry of the chessboard.

Let Q_i denote a Boolean variable indicating whether a queen is placed on square i , assuming squares are indexed from 1 to n^2 in row-major order. In addition, for every row, column, diagonal, and anti-diagonal ℓ , we introduce a Boolean variable L_ℓ , indicating whether line ℓ is active, i.e., whether it contains at least one queen. Furthermore, for each square i , let $\mathcal{L}(i)$ denote the set of the four lines passing through square i .

*Implementation available at <https://github.com/TahaRostami/Gamma/tree/gamma-plus>

²<https://github.com/curtisbright/cadical-exhaust>

³<https://github.com/curtisbright/drat-trim-t>

The domination constraints are expressed using the line variables. Indeed, every square must be dominated, and a square i is dominated if at least one of its incident lines is active. Thus, for every square i , we add the clause

$$\bigvee_{\ell \in \mathcal{L}(i)} L_\ell.$$

Next, if a line ℓ is active, then by definition at least one queen must lie on that line. Therefore, for every line ℓ , we add

$$\neg L_\ell \vee \bigvee_{i \in \ell} Q_i.$$

Since the objective is to determine whether there exists a dominating set of size at most γ , the encoding includes a cardinality constraint limiting the number of queens:

$$\sum_{i=1}^{n^2} Q_i \leq \gamma.$$

This constraint is encoded using a totalizer encoding. As observed in prior work [?], the ordering of literals significantly impacts solver performance. To preserve geometric locality, we order the queen variables according to a Hilbert curve traversal of the chessboard, following the approach introduced in previous work [1]. This ordering better reflects the board structure and improves propagation behavior in practice.

In addition to introducing the line variables, a major source of the performance improvement comes from imposing a second cardinality constraint on these variables. Since each queen lies on exactly one row, one column, one diagonal, and one anti-diagonal, a queen can activate at most four lines. Consequently, any solution using at most γ queens can activate at most 4γ lines. We therefore impose the constraint

$$\sum_{\ell \in \mathcal{L}} L_\ell \leq 4\gamma.$$

For this constraint, the line variables are ordered by decreasing line length before being passed to the mTotalizer cardinality encoder.

Symmetry Breaking

To reduce redundant exploration of symmetric solutions, we incorporate static symmetry-breaking constraints, similarly to previous work [1]. Our formulation uses a lexicographic symmetry-breaking scheme based on the encoding proposed by Harvey.

Apart from the identity transformation, the $n \times n$ chessboard has seven geometric symmetries: rotations by 90° , 180° , and 270° , as well as reflections across the horizontal axis, vertical axis, main diagonal, and anti-diagonal. Let X denote the Boolean vector associated with the original ordering of the queen variables. For each non-identity symmetry, we construct a transformed vector Y and enforce the lexicographic condition

$$X \leq Y.$$

To encode the lexicographic comparison

$$[x_1, \dots, x_N] \leq [y_1, \dots, y_N],$$

we introduce auxiliary variables $a_0, a_1, \dots, a_N$, where $N = n^2$. Intuitively, each variable a_i captures whether the suffix

$$[x_{i+1}, \dots, x_N]$$

is lexicographically smaller than or equal to the corresponding suffix

$$[y_{i+1}, \dots, y_N].$$

Following Harvey’s construction, the encoding introduces the ternary clauses

$$\begin{aligned} a_{i+1} \vee y_{i+1} \vee \neg a_i, \\ a_{i+1} \vee \neg x_{i+1} \vee \neg a_i, \end{aligned}$$

and

$$y_{i+1} \vee \neg x_{i+1} \vee \neg a_i,$$

for every $0 \leq i < N$, together with unit clauses enforcing a_0 and a_N to be true. These constraints are generated for all seven nontrivial symmetries of the chessboard. In contrast to previous work, we additionally observed that arranging the symmetry-breaking vectors according to a Hilbert curve ordering further improves solver performance in practice.

Cube-and-Conquer

To solve larger instances more effectively, we use the Cube-and-Conquer paradigm, which partitions the original SAT instance into many smaller subproblems, called cubes, that can then be solved independently by a conflict-driven SAT solver. For smaller instances, Cube-and-Conquer is often optional and may provide only moderate speedups. For larger instances, however, partitioning becomes essential, since solving the original instance directly becomes increasingly difficult.

Previous work [1] employed a Cube-and-Conquer approach based on `march_cu`. Although effective on smaller instances, that framework struggled to scale on larger queen domination instances. In the present work, we instead adopt a more scalable Cube-and-Conquer framework [2], which is designed specifically for totalizer-based encodings. For this reason, we encode the queen cardinality constraint using a standard totalizer rather than the `mTotalizer` encoding used previously.

In particular, our cubing strategy selects auxiliary variables from the totalizer encoding of the queen cardinality constraint. These variables encode partial counting information, namely whether at least a certain number of queen variables are assigned true within a sub-block of the ordered variables. Branching on such variables partitions the search space according to meaningful cardinality structure rather than only individual queen placements, producing independent subproblems that can then be solved separately. Interested readers can find a detailed explanation of this procedure in [2].

References

- [1] Rostami, T., and Bright, C. Queen Domination by SAT Solving. arXiv:2508.11945, 2025.
- [2] Battleman, Z., Reeves, J. E., and Heule, M. J. H. Problem partitioning via proof prefixes. In *Proceedings of the 28th International Conference on Theory and Applications of Satisfiability Testing (SAT 2025)*, pages 3:1–3:18. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2025.